

Loan Approval Prediction - Model Development

Step 1: Import Necessary Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import classification_report, accuracy_score, precision_score, recall_score, f1_score
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import cross_val_score
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.preprocessing import StandardScaler
from imblearn.under_sampling import RandomUnderSampler
```

Step 2: Import Cleaned Loan Dataset

```
In [2]: df = pd.read_csv(r'prosperLoanDataCleaned.csv')
df.head()
```

```
Out[2]:
```

	Term	MonthsOfEmploymentExperience	IsHomeowner	OpenCreditLines	TotalInquiries	AvailableBankcardCredit	DebtToIncomeRatio
0	36	2.000000	1	4.000000	3	1500.000000	0.250000
1	36	96.071582	0	9.260164	1	11210.225447	0.150000
2	36	19.000000	0	2.000000	5	2580.000000	0.200000
3	36	1.000000	0	7.000000	4	3626.000000	0.180000
4	36	121.000000	1	9.000000	1	178.000000	0.120000

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55106 entries, 0 to 55105
Data columns (total 17 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Term                                55106 non-null  int64
 1   MonthsOfEmploymentExperience        55106 non-null  float64
 2   IsHomeowner                        55106 non-null  int64
 3   OpenCreditLines                    55106 non-null  float64
 4   TotalInquiries                      55106 non-null  int64
 5   AvailableBankcardCredit             55106 non-null  float64
 6   DebtToIncomeRatio                  55106 non-null  float64
 7   IncomeVerifiable                   55106 non-null  int64
 8   StatedMonthlyIncome                55106 non-null  float64
 9   LoanNumber                         55106 non-null  int64
10   LoanAmount                         55106 non-null  int64
11   MonthlyInstallment                 55106 non-null  float64
12   InterestRate                       55106 non-null  float64
13   IsEmployed                         55106 non-null  int64
14   AverageCreditScore                 55106 non-null  int64
15   AnyDelinquencies                   55106 non-null  int64
16   GoodLoan                           55106 non-null  int64
dtypes: float64(7), int64(10)
memory usage: 7.1 MB
```

Step 3: Data Pre-processing

```
In [4]: df = df.drop(['LoanNumber'], axis=1)
categorical = ['Term', 'IsHomeowner', 'IsEmployed', 'AnyDelinquencies', 'IncomeVerifiable']
numerical = ['LoanAmount', 'InterestRate', 'MonthlyInstallment', 'MonthsOfEmploymentExperience', 'AverageCr
            ', 'OpenCreditLines', 'TotalInquiries', 'AvailableBankcardCredit', 'DebtToIncomeRatio', 'StatedMo
```

Handling Outliers

```
In [5]: df[numerical].skew()
```

```
Out[5]: LoanAmount                1.646822
InterestRate                0.115913
MonthlyInstallment          1.953124
MonthsOfEmploymentExperience 1.834853
AverageCreditScore         -1.351894
OpenCreditLines             1.131160
TotalInquiries              6.574083
AvailableBankcardCredit     7.593415
DebtToIncomeRatio           12.416432
StatedMonthlyIncome         44.534218
dtype: float64
```

```
In [6]: Q1 = df[numerical].quantile(0.25)
Q3 = df[numerical].quantile(0.75)
IQR = Q3 - Q1

# Calculate median
median = df[numerical].median()

# Identify outliers
outliers_lower = (df[numerical] < (Q1 - 1.5 * IQR))
outliers_upper = (df[numerical] > (Q3 + 1.5 * IQR))

# Replace outliers with median
df[numerical] = df[numerical].mask(outliers_lower, median, axis=1)
df[numerical] = df[numerical].mask(outliers_upper, median, axis=1)
```

```
In [7]: df[numerical].skew()
```

```
Out[7]: LoanAmount          1.106712
InterestRate          0.115913
MonthlyInstallment    0.980513
MonthsOfEmploymentExperience 0.856838
AverageCreditScore   -0.158725
OpenCreditLines       0.453486
TotalInquiries        1.077752
AvailableBankcardCredit 1.292952
DebtToIncomeRatio     0.428901
StatedMonthlyIncome   0.591492
dtype: float64
```

Step 4: Prepare Data for Model Training

Shuffle the Rows of a DataFrame

```
In [8]: df = df.sample(frac=1, random_state=42)
df
```

```
Out[8]:
```

	Term	MonthsOfEmploymentExperience	IsHomeowner	OpenCreditLines	TotalInquiries	AvailableBankcardCredit	De
2098	36	60.000000	1	8.000000	5	362.000000	
43905	36	9.000000	0	5.000000	9	0.000000	
46310	36	96.071582	0	9.260164	1	11210.225447	
53388	36	63.000000	1	6.000000	5	1913.000000	
32128	36	11.000000	1	9.000000	8	138.000000	
...
44732	36	49.000000	1	4.000000	5	7046.000000	
54343	36	22.000000	0	2.000000	6	0.000000	
38158	36	87.000000	0	6.000000	13	491.000000	
860	36	18.000000	1	8.000000	14	7252.000000	
15795	36	16.000000	0	6.000000	2	4368.500000	

55106 rows × 16 columns

Separate the Input Features (X) from the Target Variable (y).

```
In [9]: X = df.drop(columns=['GoodLoan'])
y = df['GoodLoan']
```

Standardize Numerical Features using Feature Scaling.

```
In [10]: StandardScaler = StandardScaler()
X[numerical] = StandardScaler.fit_transform(X[numerical])

print("Scaled Features (x):")
print(X.head())

print("\nTarget Variable (y):")
print(y.head())
```

Scaled Features (x):

	Term	MonthsOfEmploymentExperience	IsHomeowner	OpenCreditLines	\
2098	36	-0.124567	1	-0.023026	
43905	36	-1.124880	0	-0.762092	
46310	36	0.582939	0	0.287422	
53388	36	-0.065725	1	-0.515737	
32128	36	-1.085652	1	0.223329	
	TotalInquiries	AvailableBankcardCredit	DebtToIncomeRatio	\	
2098	-0.114880	-0.875305	2.405529		
43905	0.820833	-0.934278	2.680542		
46310	-1.050592	0.891976	-0.436274		
53388	-0.114880	-0.622632	1.488818		
32128	0.586905	-0.911796	0.534950		
	IncomeVerifiable	StatedMonthlyIncome	LoanAmount	MonthlyInstallment	\
2098	1	2.068814	-0.939101	-1.510281	
43905	1	-1.551569	-1.204407	-1.168564	
46310	1	-1.522892	-0.275839	-0.085855	
53388	1	-0.524802	-0.408491	-0.274842	
32128	0	-1.911839	-0.673796	-0.693528	
	InterestRate	IsEmployed	AverageCreditScore	AnyDelinquencies	
2098	1.200108	1	-0.268925	0	
43905	1.817964	1	-0.858709	0	
46310	1.079491	1	0.025967	1	
53388	0.848102	1	-1.153601	0	
32128	-0.301455	1	-2.038278	0	

Target Variable (y):

2098 1
 43905 0
 46310 0
 53388 0
 32128 1

Name: GoodLoan, dtype: int64

Split the Data into Train Test and Validation Sets.

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.25, random_state=42)
```

Perform Feature Selection

```
In [12]: # Instantiate SelectKBest with ANOVA F-value as the score function
k_best_selector = SelectKBest(score_func=f_classif, k=12) # Select top 12 features

# Fit SelectKBest to training data
X_train_selected = k_best_selector.fit_transform(X_train, y_train)

# Print the names of the selected features
selected_feature_indices = k_best_selector.get_support(indices=True)
selected_feature_names = X.columns[selected_feature_indices]
print("Selected Features:")
for feature in selected_feature_names:
    print(feature)

X_train = X_train[selected_feature_names]
X_val = X_val[selected_feature_names]
X_test = X_test[selected_feature_names]
```

Selected Features:
 Term
 IsHomeowner
 OpenCreditLines
 TotalInquiries
 AvailableBankcardCredit
 DebtToIncomeRatio
 IncomeVerifiable
 StatedMonthlyIncome
 MonthlyInstallment
 InterestRate
 AverageCreditScore
 AnyDelinquencies

Step 6: Model Building

Subset Selection

```
In [13]: def forward_stepwise_selection(model, X_train, X_test, y_train, y_test):
    remaining_features = list(X_train.columns)
    selected_features = []
    prev_score = 0
    while True:
        best_score = 0
        for feature in remaining_features:
            model.fit(X_train[selected_features + [feature]], y_train)
            y_pred = model.predict(X_test[selected_features + [feature]])
            #score = f1_score(y_test, y_pred, average='weighted')
            fpr, tpr, _ = roc_curve(y_test, y_pred)
            score = auc(fpr, tpr)
            #score = accuracy_score(y_test, y_pred)
            if score > best_score:
                best_score = score
                best_feature = feature
        if best_score > prev_score:
            prev_score = best_score
            selected_features.append(best_feature)
            remaining_features.remove(best_feature)
        else:
            break
    return selected_features
```

Logistic Regression

```
In [14]: LR = LogisticRegression(max_iter=1000, class_weight={0:3, 1:1}, random_state=42)

# Best Subset Selection
# best_subset_LR = forward_stepwise_selection(LR, X_train, X_val, y_train, y_val)
# print("\nSelected Features:", best_subset_LR)

# Hyper-parameter Tuning
# param_grid = {
#     'C': [0.001, 0.01, 0.1, 1, 10, 100]
# }
# grid_search = GridSearchCV(LR, param_grid, cv=5, scoring='roc_auc')
# grid_search.fit(X_train[best_subset_LR], y_train)
# best_params_LR = grid_search.best_params_
# print("\nBest Parameters:", best_params_LR)

best_subset_LR = ['InterestRate', 'Term', 'TotalInquiries', 'StatedMonthlyIncome', 'AnyDelinquencies',
                  'DebtToIncomeRatio', 'IsHomeowner', 'AverageCreditScore', 'MonthlyInstallment', 'IncomeVerifiable']
print("\nSelected Features:", best_subset_LR)
best_params_LR = {'C': 1}
print("\nBest Parameters:", best_params_LR)

# Model with best parameters
LR = LogisticRegression(**best_params_LR, max_iter=1000, class_weight={0:3, 1:1}, random_state=42)
```

```

# Perform cross-validation
cv_scores = cross_val_score(LR, X_train[best_subset_LR], y_train, cv=5)

# Print cross-validation scores
print("\nCross-validation scores:", cv_scores)

# Calculate mean validation accuracy
validation_acc_LR = np.mean(cv_scores)
print("\nValidation accuracy:", np.mean(validation_acc_LR))

# Fit the model with selected features
LR.fit(X_train[best_subset_LR], y_train)

# Calculate training accuracy
y_pred = LR.predict(X_train[best_subset_LR])
train_acc_LR = accuracy_score(y_train, y_pred)
print("Training accuracy:", train_acc_LR)

# Calculate testing accuracy
y_pred = LR.predict(X_test[best_subset_LR])
test_acc_LR = accuracy_score(y_test, y_pred)
print("Testing accuracy:", test_acc_LR)

# Calculate precision, recall, and F1 score
precision_LR = precision_score(y_test, y_pred, average='weighted')
recall_LR = recall_score(y_test, y_pred, average='weighted')
f1_LR = f1_score(y_test, y_pred, average='weighted')

# Build a classification report
classification_rep_LR = classification_report(y_test, y_pred)

# Plot a confusion matrix
confusion_matrix_LR = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix_LR, annot=True, fmt='d', cmap='Oranges')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Print the classification report
print("Classification Report:\n", classification_rep_LR)
print("Precision:", precision_LR)
print("Recall:", recall_LR)
print("F1 Score:", f1_LR)

print("\nThe accuracy of the Logistic Regression Model is:", round(test_acc_LR*100,2),"% .")

fpr_LR, tpr_LR, _ = roc_curve(y_test, y_pred)
auc_LR = auc(fpr_LR, tpr_LR)

print("\nThe ROC-AUC Score of the Logistic Regression Model is:", auc_LR, ".")

```

Selected Features: ['InterestRate', 'Term', 'TotalInquiries', 'StatedMonthlyIncome', 'AnyDelinquencies', 'DebtToIncomeRatio', 'IsHomeowner', 'AverageCreditScore', 'MonthlyInstallment', 'IncomeVerifiable']


Best Parameters: {'C': 1}

Cross-validation scores: [0.57034221 0.57673695 0.57414449 0.57569997 0.57639129]

Validation accuracy: 0.5746629796059454

Training accuracy: 0.5752160387141376

Testing accuracy: 0.577425598838616



Classification Report:

	precision	recall	f1-score	support
0	0.43	0.83	0.57	5481
1	0.84	0.45	0.59	11051
accuracy			0.58	16532
macro avg	0.64	0.64	0.58	16532
weighted avg	0.71	0.58	0.58	16532

Precision: 0.7051338751472839

Recall: 0.577425598838616

F1 Score: 0.5810628014428235

The accuracy of the Logistic Regression Model is: 57.74 % .

The ROC-AUC Score of the Logistic Regression Model is: 0.6407912785179315 .

Extreme Gradient Boosting (XGBoost)

```
In [15]: weight=len(y_train[y_train == 0]) / len(y_train[y_train == 1])
XGB = XGBClassifier(random_state=42,scale_pos_weight=weight)

# Best Subset Selection
# best_subset_XGB = forward_stepwise_selection(XGB,X_train, X_val, y_train, y_val)

# Hyper-parameter Tuning
# param_grid = {
#     'n_estimators': [100, 200, 300], # Number of boosting rounds
#     'learning_rate': [0.05, 0.1, 0.2], # Step size shrinkage used in update to prevent overfitting
#     'max_depth': [3, 4, 5], # Maximum depth of a tree, deeper trees can model more complex relationships
# }

# grid_search = GridSearchCV(XGB, param_grid, cv=5, scoring='roc_auc')
# grid_search.fit(X_train[best_subset_XGB], y_train)
# best_params_XGB = grid_search.best_params_
```

```

# print("\nSelected Features:", best_subset_XGB)
# print("\nBest Parameters:", best_params_XGB)

best_subset_XGB = ['InterestRate', 'MonthlyInstallment', 'AverageCreditScore', 'Term', 'AnyDelinquencies',
'IsHomeowner', 'DebtToIncomeRatio', 'StatedMonthlyIncome', 'OpenCreditLines']
print("\nSelected Features:", best_subset_XGB)

best_params_XGB = {'learning_rate': 0.1, 'max_depth': 4, 'n_estimators': 100}
print("\nBest Parameters:", best_params_XGB)

# Model with best parameters
XGB = XGBClassifier(**best_params_XGB, random_state=42, scale_pos_weight=weight)

# Perform cross-validation
cv_scores = cross_val_score(XGB, X_train[best_subset_XGB], y_train, cv=5)

# Print cross-validation scores
print("\nCross-validation scores:", cv_scores)

# Calculate mean validation accuracy
validation_acc_XGB = np.mean(cv_scores)
print("\nValidation accuracy:", np.mean(validation_acc_XGB))

# Fit the model with selected features
XGB.fit(X_train[best_subset_XGB], y_train)

# Calculate training accuracy
y_pred = XGB.predict(X_train[best_subset_XGB])
train_acc_XGB = accuracy_score(y_train, y_pred)
print("Training accuracy:", train_acc_XGB)

# Calculate testing accuracy
y_pred = XGB.predict(X_test[best_subset_XGB])
test_acc_XGB = accuracy_score(y_test, y_pred)
print("Testing accuracy:", test_acc_XGB)

# Calculate precision, recall, and F1 score
precision_XGB = precision_score(y_test, y_pred, average='weighted')
recall_XGB = recall_score(y_test, y_pred, average='weighted')
f1_XGB = f1_score(y_test, y_pred, average='weighted')

# Build a classification report
classification_rep_XGB = classification_report(y_test, y_pred)

# Plot a confusion matrix
confusion_matrix_XGB = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix_XGB, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Print the classification report
print("Classification Report:\n", classification_rep_XGB)
print("Precision:", precision_XGB)
print("Recall:", recall_XGB)
print("F1 Score:", f1_XGB)

print("\nThe accuracy of the Extreme Gradient Boosting Model is:", round(test_acc_XGB*100,2), "% .")

fpr_XGB, tpr_XGB, _ = roc_curve(y_test, y_pred)
auc_XGB = auc(fpr_XGB, tpr_XGB)

print("\nThe ROC-AUC Score of the Extreme Gradient Boosting (XGBoost) is:", auc_XGB, ".")

```


7/26/24, 12:37 PM

Loan Approval Prediction - Model Development

Selected Features: ['InterestRate', 'MonthlyInstallment', 'AverageCreditScore', 'Term', 'AnyDelinquencies', 'IsHomeowner', 'DebtToIncomeRatio', 'StatedMonthlyIncome', 'OpenCreditLines']


Best Parameters: {'learning_rate': 0.1, 'max_depth': 4, 'n_estimators': 100}

Cross-validation scores: [0.63480816 0.64379537 0.6389561 0.63878327 0.64362254]

Validation accuracy: 0.6399930867611475

Training accuracy: 0.652886277220878

Testing accuracy: 0.6408782966368255



Classification Report:

	precision	recall	f1-score	support
0	0.47	0.73	0.57	5481
1	0.82	0.60	0.69	11051
accuracy			0.64	16532
macro avg	0.64	0.66	0.63	16532
weighted avg	0.70	0.64	0.65	16532

Precision: 0.7028706412690978
Recall: 0.6408782966368256
F1 Score: 0.6512743583745236

The accuracy of the Extreme Gradient Boosting Model is: 64.09 % .

The ROC-AUC Score of the Extreme Gradient Boosting (XGBoost) is: 0.6634241410232973 .

Under Sampling the Dataset to handle the Imbalance in Data

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.25, random_state=42)

undersampler = RandomUnderSampler(random_state=42)
X_train, y_train = undersampler.fit_resample(X_train, y_train)
```

K - Nearest Neighbour (KNN)

```
In [17]: KNN = KNeighborsClassifier()

# Best Subset Selection
# best_subset_KNN = forward_stepwise_selection(KNN,X_train, X_val, y_train, y_val)
# print("\nSelected Features:", best_subset_KNN)

# Hyper-parameter Tuning
# param_grid = {
#     'n_neighbors': [3, 5, 7, 9, 11], # Different values of k
#     'weights': ['uniform', 'distance'], # Weighting scheme
#     'metric': ['euclidean', 'manhattan'] # Distance metric
# }
# grid_search = GridSearchCV(KNN, param_grid, cv=5, scoring='roc_auc')
# grid_search.fit(X_train[best_subset_KNN], y_train)
# best_params_KNN = grid_search.best_params_
# print("\nBest Parameters:", best_params_KNN)

best_subset_KNN = ['MonthlyInstallment', 'LoanAmount', 'Term', 'AverageCreditScore', 'AnyDelinquencies',
print("\nSelected Features:", best_subset_KNN)
best_params_KNN = {'metric': 'manhattan', 'n_neighbors': 11, 'weights': 'uniform'}
print("\nBest Parameters:", best_params_KNN)

# Model with best parameters
KNN = KNeighborsClassifier(**best_params_KNN)

# Perform cross-validation
cv_scores = cross_val_score(KNN, X_train[best_subset_KNN], y_train, cv=5)

# Print cross-validation scores
print("\nCross-validation scores:", cv_scores)

# Calculate mean validation accuracy
validation_acc_KNN = np.mean(cv_scores)
print("\nValidation accuracy:", np.mean(validation_acc_KNN))

# Fit the model with selected features
KNN.fit(X_train[best_subset_KNN], y_train)

# Calculate training accuracy
y_pred = KNN.predict(X_train[best_subset_KNN])
train_acc_KNN = accuracy_score(y_train, y_pred)
print("Training accuracy:", train_acc_KNN)

# Calculate testing accuracy
y_pred = KNN.predict(X_test[best_subset_KNN])
test_acc_KNN = accuracy_score(y_test, y_pred)
print("Testing accuracy:", test_acc_KNN)

# Calculate precision, recall, and F1 score
precision_KNN = precision_score(y_test, y_pred, average='weighted')
recall_KNN = recall_score(y_test, y_pred, average='weighted')
f1_KNN = f1_score(y_test, y_pred, average='weighted')

# Build a classification report
classification_rep_KNN = classification_report(y_test, y_pred)

# Plot a confusion matrix
confusion_matrix_KNN = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
```

```

sns.heatmap(confusion_matrix_KNN, annot=True, fmt='d', cmap='Greens')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Print the classification report
print("Classification Report:\n", classification_rep_KNN)
print("Precision:", precision_KNN)
print("Recall:", recall_KNN)
print("F1 Score:", f1_KNN)

print("\nThe accuracy of the K - Nearest Neighbour Model is:", round(test_acc_KNN*100,2),"% .")

fpr_KNN, tpr_KNN, _ = roc_curve(y_test, y_pred)
auc_KNN = auc(fpr_KNN, tpr_KNN)

print("\nThe ROC-AUC Score of the K - Nearest Neighbour Model is:", auc_KNN, ".")

```

Selected Features: ['MonthlyInstallment', 'LoanAmount', 'Term', 'AverageCreditScore', 'AnyDelinquencies', 'InterestRate']


Best Parameters: {'metric': 'manhattan', 'n_neighbors': 11, 'weights': 'uniform'}

Cross-validation scores: [0.61836969 0.63467279 0.62962113 0.61859931 0.63206247]

Validation accuracy: 0.6266650774958114

Training accuracy: 0.6920180031229908

Testing accuracy: 0.6130466340047178



Classification Report:

	precision	recall	f1-score	support
0	0.45	0.70	0.55	3695
1	0.79	0.57	0.66	7327
accuracy			0.61	11022
macro avg	0.62	0.63	0.60	11022
weighted avg	0.68	0.61	0.62	11022

Precision: 0.6750960859798549

Recall: 0.6130466340047178

F1 Score: 0.6236435188070928

The accuracy of the K - Nearest Neighbour Model is: 61.3 % .

The ROC-AUC Score of the K - Nearest Neighbour Model is: 0.6336925191697419 .

Multilayer Perceptron (MLP) Classifier

```
In [18]: MLP = MLPClassifier(random_state=42, max_iter=1000)

# Best Subset Selection
# best_subset_MLP = forward_stepwise_selection(MLP,X_train, X_val, y_train, y_val)
best_subset_MLP = ['InterestRate', 'MonthlyInstallment', 'LoanAmount', 'TotalInquiries',
                   'StatedMonthlyIncome', 'AnyDelinquencies', 'AverageCreditScore']
print("\nSelected Features:", best_subset_MLP)

# Hyper-parameter Tuning
# param_grid = {
#     'hidden_layer_sizes': [(50,), (100,), (50, 50), (100, 50)], # Different architectures for hidden
#     'activation': ['relu', 'tanh'], # Activation functions
#     'alpha': [0.0001, 0.001, 0.01], # L2 regularization parameter
# }
# grid_search = GridSearchCV(MLP, param_grid, cv=5, scoring='roc_auc')
# grid_search.fit(X_train[best_subset_MLP], y_train)
# best_params_MLP = grid_search.best_params_
best_params_MLP = {'activation': 'relu', 'alpha': 0.01, 'hidden_layer_sizes': (50,)}
print("\nBest Parameters:", best_params_MLP)

# Model with best parameters
MLP = MLPClassifier(**best_params_MLP, random_state=42, max_iter=1000)

# Perform cross-validation
cv_scores = cross_val_score(MLP, X_train[best_subset_MLP], y_train, cv=5)

# Print cross-validation scores
print("\nCross-validation scores:", cv_scores)

# Calculate mean validation accuracy
validation_acc_MLP = np.mean(cv_scores)
print("\nValidation accuracy:", np.mean(validation_acc_MLP))

# Fit the model with selected features
MLP.fit(X_train[best_subset_MLP], y_train)

# Calculate training accuracy
y_pred = MLP.predict(X_train[best_subset_MLP])
train_acc_MLP = accuracy_score(y_train, y_pred)
print("Training accuracy:", train_acc_MLP)

# Calculate testing accuracy
y_pred = MLP.predict(X_test[best_subset_MLP])
test_acc_MLP = accuracy_score(y_test, y_pred)
print("Testing accuracy:", test_acc_MLP)

# Calculate precision, recall, and F1 score
precision_MLP = precision_score(y_test, y_pred, average='weighted')
```

```
recall_MLP = recall_score(y_test, y_pred, average='weighted')
f1_MLP = f1_score(y_test, y_pred, average='weighted')

# Build a classification report
classification_rep_MLP = classification_report(y_test, y_pred)

# Plot a confusion matrix
confusion_matrix_MLP = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix_MLP, annot=True, fmt='d', cmap='Purples')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Print the classification report
print("Classification Report:\n", classification_rep_MLP)
print("Precision:", precision_MLP)
print("Recall:", recall_MLP)
print("F1 Score:", f1_MLP)

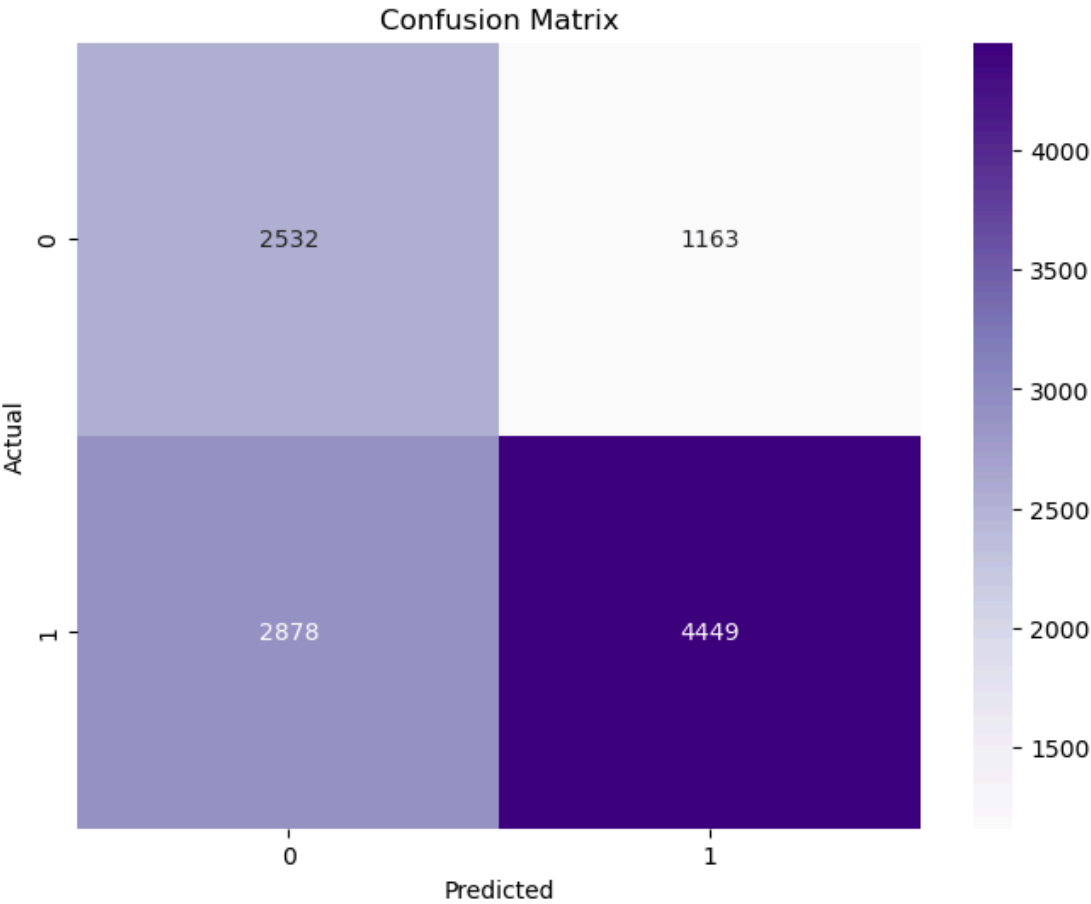
print("\nThe accuracy of the Multilayer Perceptron (MLP) Classifier Model is:", round(test_acc_MLP*100,2))

fpr_MLP, tpr_MLP, _ = roc_curve(y_test, y_pred)
auc_MLP = auc(fpr_MLP, tpr_MLP)

print("\nThe ROC-AUC Score of the Multilayer Perceptron (MLP) Classifier Model is:",auc_MLP,".")
```

Selected Features: ['InterestRate', 'MonthlyInstallment', 'LoanAmount', 'TotalInquiries', 'StatedMonthlyIncome', 'AnyDelinquencies', 'AverageCreditScore']

```
Validation accuracy: 0.6496741689946086
Training accuracy: 0.662579222926426
Testing accuracy: 0.6333696243875885
```



Classification Report:

	precision	recall	f1-score	support
0	0.47	0.69	0.56	3695
1	0.79	0.61	0.69	7327
accuracy			0.63	11022
macro avg	0.63	0.65	0.62	11022
weighted avg	0.68	0.63	0.64	11022

Precision: 0.6838990016433403
Recall: 0.6333696243875885
F1 Score: 0.6436009997482628

The accuracy of the Multilayer Perceptron (MLP) Classifier Model is: 63.34 % .

The ROC-AUC Score of the Multilayer Perceptron (MLP) Classifier Model is: 0.6462282809258506 .

Support Vector Machine (SVM)

```
In [19]: SVM = SVC(random_state=42)

# Best Subset Selection
# best_subset_SVM = forward_stepwise_selection(SVM,X_train, X_val, y_train, y_val)
best_subset_SVM = ['InterestRate', 'MonthlyInstallment', 'TotalInquiries', 'AvailableBankcardCredit',
'AverageCreditScore', 'StatedMonthlyIncome', 'LoanAmount', 'AnyDelinquencies', 'IsHomeowner', 'IncomeVe
print("\nSelected Features:", best_subset_SVM)

# Hyper-parameter Tuning
# param_grid = {
#     'C': [0.1, 1, 10, 100], # Regularization parameter
#     'kernel': ['linear', 'rbf', 'poly'],# Kernel type
#     'gamma': ['scale', 'auto'], # Kernel coefficient for RBF and poly kernels
#     'degree': [2, 3, 4], # Degree for poly kernel
#     'coef0': [0.0, 0.1, 0.5] # Independent term in poly and sigmoid kernels
# }
```

```

# random_search = RandomizedSearchCV(SVM, param_grid, cv=5, scoring='roc_auc')
# random_search.fit(X_train[best_subset_SVM], y_train)
# best_params_SVM = random_search.best_params_
best_params_SVM = {'kernel': 'rbf', 'gamma': 'auto', 'degree': 3, 'coef0': 0.0, 'C': 1}
print("\nBest Parameters:", best_params_SVM)

# Model with best parameters
SVM = SVC(**best_params_SVM)

# Perform cross-validation
cv_scores = cross_val_score(SVM, X_train[best_subset_SVM], y_train, cv=5)

# Print cross-validation scores
print("\nCross-validation scores:", cv_scores)

# Calculate mean validation accuracy
validation_acc_SVM = np.mean(cv_scores)
print("\nValidation accuracy:", np.mean(validation_acc_SVM))

# Fit the model with selected features
SVM.fit(X_train[best_subset_SVM], y_train)

# Calculate training accuracy
y_pred = SVM.predict(X_train[best_subset_SVM])
train_acc_SVM = accuracy_score(y_train, y_pred)
print("Training accuracy:", train_acc_SVM)

# Calculate testing accuracy
y_pred = SVM.predict(X_test[best_subset_SVM])
test_acc_SVM = accuracy_score(y_test, y_pred)
print("Testing accuracy:", test_acc_SVM)

# Calculate precision, recall, and F1 score
precision_SVM = precision_score(y_test, y_pred, average='weighted')
recall_SVM = recall_score(y_test, y_pred, average='weighted')
f1_SVM = f1_score(y_test, y_pred, average='weighted')

# Build a classification report
classification_rep_SVM = classification_report(y_test, y_pred)

# Plot a confusion matrix
confusion_matrix_SVM = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix_SVM, annot=True, fmt='d', cmap='Greys')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Print the classification report
print("Classification Report:\n", classification_rep_SVM)
print("Precision:", precision_SVM)
print("Recall:", recall_SVM)
print("F1 Score:", f1_SVM)

print("\nThe accuracy of the Support Vector Machine Model is:", round(test_acc_SVM*100,2),"% .")

fpr_SVM, tpr_SVM, _ = roc_curve(y_test, y_pred)
auc_SVM = auc(fpr_SVM, tpr_SVM)

print("\nThe ROC-AUC Score of the Support Vector Machine Model is:", auc_SVM, ".")

```


7/26/24, 12:37 PM

Loan Approval Prediction - Model Development

Selected Features: ['InterestRate', 'MonthlyInstallment', 'TotalInquiries', 'AvailableBankcardCredit', 'AverageCreditScore', 'StatedMonthlyIncome', 'LoanAmount', 'AnyDelinquencies', 'IsHomeowner', 'IncomeVerifiable']


Best Parameters: {'kernel': 'rbf', 'gamma': 'auto', 'degree': 3, 'coef0': 0.0, 'C': 1}

Cross-validation scores: [0.65258324 0.65350172 0.652124 0.64546498 0.64813964]

Validation accuracy: 0.650362715942214

Training accuracy: 0.6634058969413061

Testing accuracy: 0.6199419343131918



Classification Report:

	precision	recall	f1-score	support
0	0.46	0.72	0.56	3695
1	0.80	0.57	0.67	7327
accuracy			0.62	11022
macro avg	0.63	0.65	0.61	11022
weighted avg	0.69	0.62	0.63	11022

Precision: 0.6868495709927309
Recall: 0.6199419343131918
F1 Score: 0.6301194637007504

The accuracy of the Support Vector Machine Model is: 61.99 % .

The ROC-AUC Score of the Support Vector Machine Model is: 0.6453853090862888 .

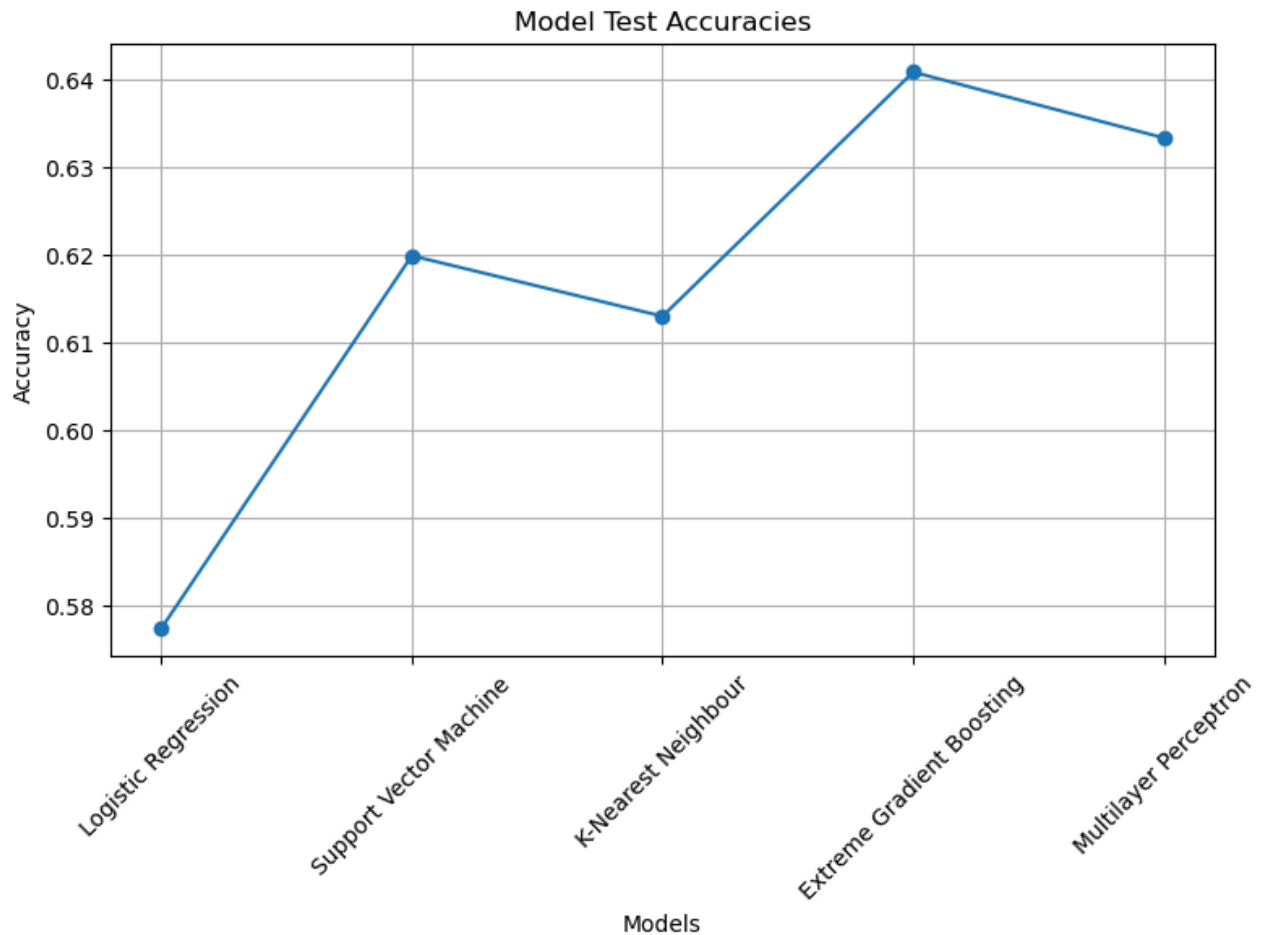
Step 7: Model Comparison

```
In [20]: models = ['Logistic Regression', 'Support Vector Machine', 'K-Nearest Neighbour', 'Extreme Gradient Boos
accuracies = [test_acc_LR, test_acc_SVM ,test_acc_KNN, test_acc_XGB, test_acc_MLP] # Example accuracies

plt.figure(figsize=(8, 6))
plt.plot(models, accuracies, marker='o')
```

```
plt.title('Model Test Accuracies')
plt.xlabel('Models')
plt.ylabel('Accuracy')

plt.grid(True)
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.tight_layout() # Adjust layout to prevent cropping of labels
plt.show()
```



```
In [21]: import matplotlib.pyplot as plt

models = ['Logistic Regression', 'Support Vector Machine', 'K-Nearest Neighbour', 'Extreme Gradient Boos

train accuracies = [train_acc_LR, train_acc_SVM, train_acc_KNN, train_acc_XGB, train_acc_MLP]
validation accuracies = [validation_acc_LR, validation_acc_SVM, validation_acc_KNN, validation_acc_XGB,
test accuracies = [test_acc_LR, test_acc_SVM, test_acc_KNN, test_acc_XGB, test_acc_MLP]
f1_scores = [f1_LR, f1_SVM, f1_KNN, f1_XGB, f1_MLP]

plt.figure(figsize=(10, 6))


# Plot train accuracies
plt.plot(models, train_accuracies, marker='o', label='Train')

# Plot validation accuracies
plt.plot(models, validation_accuracies, marker='o', label='Validation')

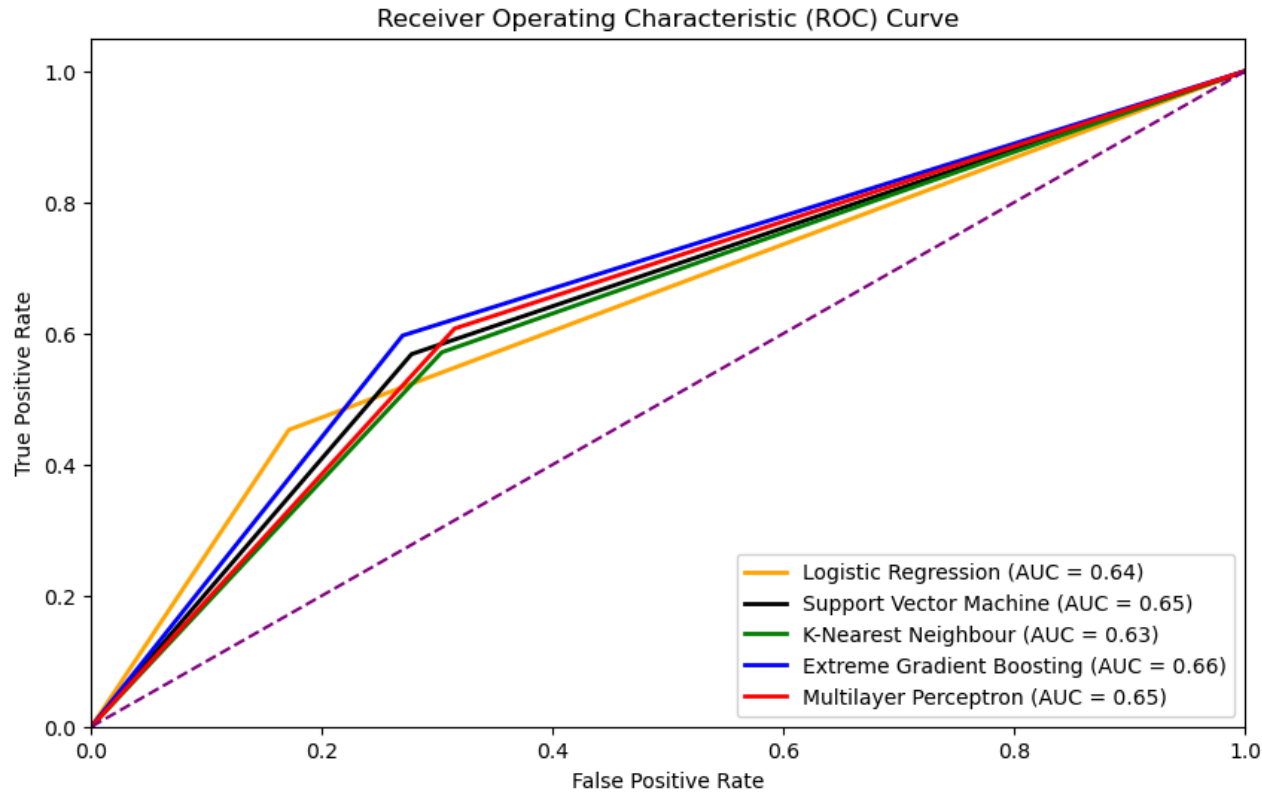
# Plot test accuracies
plt.plot(models, test_accuracies, marker='o', label='Test')

plt.title('Model Accuracies')
plt.xlabel('Models')
plt.ylabel('Accuracy')
plt.grid(True)
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
```

```
plt.legend() # Show Legend
plt.tight_layout() # Adjust layout to prevent cropping of labels
plt.show()
```



```
In [22]: plt.figure(figsize=(10, 6))
plt.plot(fpr_LR, tpr_LR, color='orange', lw=2, label=f'Logistic Regression (AUC = {auc_LR:.2f})')
plt.plot(fpr_SVM, tpr_SVM, color='black', lw=2, label=f'Support Vector Machine (AUC = {auc_SVM:.2f})')
plt.plot(fpr_KNN, tpr_KNN, color='green', lw=2, label=f'K-Nearest Neighbour (AUC = {auc_KNN:.2f})')
plt.plot(fpr_XGB, tpr_XGB, color='blue', lw=2, label=f'Extreme Gradient Boosting (AUC = {auc_XGB:.2f})')
plt.plot(fpr_MLP, tpr_MLP, color='red', lw=2, label=f'Multilayer Perceptron (AUC = {auc_MLP:.2f})')
plt.plot([0, 1], [0, 1], color='purple', linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



```
In [24]: roc_auc_scores = [auc_LR,auc_SVM,auc_KNN,auc_XGB,auc_MLP]
df_table = pd.DataFrame({
    'Model': models,
    'Train Accuracy': train_accuracies,
    'Validation Accuracy': validation_accuracies,
    'Test Accuracy': test_accuracies,
    'ROC-AUC Score': roc_auc_scores,
    'F1-Score': f1_scores
})

# Print the DataFrame
df_table
```

Out[24]:

	Model	Train Accuracy	Validation Accuracy	Test Accuracy	ROC-AUC Score	F1-Score
0	Logistic Regression	0.575216	0.574663	0.577426	0.640791	0.581063
1	Support Vector Machine	0.663406	0.650363	0.619942	0.645385	0.630119
2	K-Nearest Neighbour	0.692018	0.626665	0.613047	0.633693	0.623644
3	Extreme Gradient Boosting	0.652886	0.639993	0.640878	0.663424	0.651274
4	Multilayer Perceptron	0.662579	0.649674	0.633370	0.646228	0.643601

```
In [ ]:
In [ ]:
```